Univariate Time Series Forecasting using ES-RNN and N-BEATS

Yves Greatti
New York University
yves.greatti@nyu.edu

Zian Chen
New York University
zian.chen@nyu.edu

Sunjoo Park
New York University
sunjoo.park@nyu.edu

Abstract

In this work, we seek to replicate and improve the results reached by two neural networks: **ES-RNN** and **N-BEATS** on the M4 dataset competition. We also run different experimentations to compare the performances of these two deep learning techniques to a more classical statistical approach like Gaussian Process ($\mathcal{GP}$). We demonstrate that although Gaussian processes could be powerful for sampling tasks and simpler to configure, these neural networks outperform it for forecasting. Neural networks could have an overhead in term of the number of hyper-parameters to tune, but when using batching they scale up very easily and generalize well to a large number of time series (100K for M4). We are thus, pretty confident that, the two neural networks could forecast with the appropriate setting of hyper-parameters, other univariate time series beyond the M4 dataset.

1 Introduction

The dataset is provided by the M4 competition organized by the International Institute of Forecasters. These competitions also known as the "Makridakis competitions", have been happening since 1982, roughly every 10 years with an increasing number of time series to forecast starting from 1000 in 1982 to 100,000 in 2018, for the M4 competition. They attract people from academia as well practitioners, the last winner is Slawek Smyl, from Uber Technologies. The model used by Smyl is a hybrid approach combining Holts-Winter smoothing techniques with a recurrent neural network (RNN). Boris Oreshkin et al. want to challenge the conclusion that, mixed techniques are the future by proposing a pure DL model with interpretable outputs: **N-Beats**, which they claim outperforms **ES-RNN**, Smyl's model. N-Beats could be a very deep DL model and it is important for their authors that, in addition to hight accuracy, N-Beats has interpretable outputs which can match established statistical models such as ETS and ARIMA, robust, efficient, and automatic models.

2 Problem Description

Univariate point forecasting in discrete time consists in predicting future values given a series of observations.

- Given an history $[y_1, \cdots, y_T]$
- TASK: predict $[y_{T+1}, \cdots, y_{T+H}]$
- H: horizon, T: length of observations, m: periodicity of the data
- Standard scale-free metrics in the practice of forecasting:
    - sMAPE: symmetric Mean Absolute Percentage Error
    - MASE: Mean Absolute Scaled Error

For the purpose of this research, we use the predictive accuracy metrics sMAPE and MASE which are the standard metrics used across the forecasting community, disregarding OWA which is a metric defined by Makridakis and requires $MASE_{Naive2}$ as part of its computation, which is generated by a random walk model set up by the organizers of the competition.

$$\text{sMAPE} = \frac{2}{H} \sum_{i=1}^{H} \frac{|y_{T+i} - \hat{y}_{T+i}|}{|y_{T+i}| + |\hat{y}_{T+i}|} \, , \text{MASE} = \frac{1}{H} \sum_{i=1}^{H} \frac{|y_{T+i} - \hat{y}_{T+i}|}{\frac{1}{T+H-m} \sum_{j=m+1}^{T+H} |y_j - y_{j-m}|}$$

2.1 Data Description

M4 dataset extends the previous three competitions by increasing significantly the number of series. The time series are grouped into six categories and can be split into two groups: high-frequency data (weekly, daily, hourly) along with low-frequency data (yearly, quarterly and monthly). They were built from real, multiple, and diverse sources and are divided into six data types: demographic, finance, industry, macro, micro and other. The minimum of observations for the training set are 79 for hourly, 652 for daily, 67 for weekly, 24 for monthly, 8 for quarterly and 7 for yearly series. They have different length raging from 7 to about 1000 observations.

Frequency	Demographic	Finance	Industry	Macro	Micro	Other	Total
Yearly	1,088	6,519	3,716	3,903	6,538	1,236	23,000
Quarterly	1,858	5,305	4,637	5,315	6,020	865	24,000
Monthly	5,728	10,987	10,017	10,016	10,975	277	48,000
Weekly	24	164	6	41	112	12	359
Daily	10	1,559	422	127	1,476	633	4,227
Hourly	0	0	0	0	0	414	414
Total	8,708	24,534	18,798	19,402	25,1212	3,437	100,000

Table 1: M4 data by type and series frequency

	Mean	Std-Dev	Min	25%	50%	75%	Max
Yearly	25	24	7	14	23	34	829
Quarterly	84	51	8	54	80	107	858
Monthly	198	137	24	64	184	288	2776
Weekly	1009	707	67	366	921	1590	2584
Daily	2343	1756	79	309	2926	4183	9905
Hourly	805	127	652	652	912	912	912

Table 2: M4 series length statistics

3 ES-RNN

3.1 Model Description

ES-RNN algorithm consists in two major layers: a preprocessing layer which uses Holt-Winter smoothing technic to extract trend and seasonality parameters and a dilated RNN network. The Holt-Winter parameters are part of the back-propagation and are tuned as the network learns the characteristics of the time series. The ES-RNN model uses a modified version of the Holt-Winters formula, in which there is no local linear level coefficient, and level and seasonality terms are scaled instead of subtracted. The linear forecast is replaced by an RNN, which has for input the normalized and de-seasonalized prior observations.

$$l_t = \alpha \frac{y_t}{s_{t-m}} + (1 - \alpha)l_{t-1}$$

$$s_t = \beta \frac{y_t}{l_{t-1}} + (1 - \beta)s_{t-m}$$

$$\hat{y}_{\textbf{win}} = \textbf{ES-RNN}(\frac{y_{ti}}{s_{ti}l_{ti}})$$

$$y_{\textbf{truth}} = (\frac{y_{to}}{s_{to}l_{to}})$$

where l is a state variable, s is a seasonality coefficient, α and β are network parameters, m is the periodicity of the data.

To model long-term dependencies in time series, ES-RNN uses dilated RNNs which allow to stack RNN cells by having skip connections between RNN cells . The dilation parameters are given in table Table 3. DRNN cells can be vanilla RNN, LSTM, GRU, GRU cells provides the higher accuracy.

Time Series	Dilations
Quarterly	(1, 2), (4, 8)
Monthly	(1, 3), (6, 12)
Daily	(1, 7), (14, 28)
Yearly	(1, 2), (2, 6)
Weekly	(1, 14), (14, 28)
Hourly	(1, 24), (24, 48)

Table 3: ES-RNN dilation parameters

3.2   Data Preparation

The M4 time series have different length, to allow vectorization and to take advantage of the GPUs, they are chopped using a predetermined cut off value. To the exception of the daily time series, for which the cut-off value is 200, 72 is used as the default value (Table 4).

	Filtering rates (in %)				
Hourly	**Daily**	**Weekly**	**Monthly**	**Quarterly**	**Yearly**
6.93	6.93	18.11	26.83	45.42	60.61

Table 4: Percentage of time series eliminated by the cut-off value

The batch size is by default 1024, we observe that, it affects the overall performance of the network, depending the time series frequency, too small the validation loss may jump up and down, and in some cases the larger the batch size, the better.

For an hourly frequency, the effect of the batch size is given in table Table 6 .

	sMAPE per category					
Batch size	**Demographic**	**Finance**	**Industry**	**Macro**	**Micro**	**Other**	**Total**
512	6.32	3.34	3.95	2.52	2.43	3.08	3.02
1024	6.20	3.27	3.9	2.52	2.37	3.07	2.97
2048	6.26	3.32	3.92	2.49	2.42	3.1	3.01

Table 5: ES-RNN batch size statistics

Each time series is split into two windows: a backcast and forecast window. Observations in each window are normalized and de-seasonalized using the coefficients of the exponential smoothing. In addition, a one-hot representation of the time series category is added to the input window.


Figure 1: Levels are a smoothing version of the original time series and seasonality coefficients are between 0 and 1.


Figure 2: The input and ouput of the RNN are normalized and de-seasonalized, compared to N-BEATS which are not scaled.

3.3 Deep learning architecture

ES-RNN layers from top to bottom consist in a stack of dilated RNNs followed by a fully connected layer, an activation function and a dense layer.

ES-RNN
time series $\in \mathbb{R}^{b \times l}$
RNN $\in \mathbb{R}^{l \times s}$
RNN $\in \mathbb{R}^{s \times s}$
$\cdots$
RNN $\in \mathbb{R}^{s \times s}$
dense $\in \mathbb{R}^{s \times s}$
Tanh
dense $\in \mathbb{R}^{s \times h}$

Table 6: ES-RNN architecture. Let b the batch size, l the length of the time series, s the embedding dimensional state size and h the length of the prediction.

Dropout or batch normalization did not provide any significant improvement, the smoothing and normalization of the data seemed enough for a convergence to a local minimum. In both networks, using the pinball loss function brought definitely a boost in performance compared to L2 or L1 losses. The pinball loss function is popular in forecasting. Two different slopes are used depending whether the forecast $\hat{y}_t$ value is greater or smaller than the truth value y_t. Two models are then trained, corresponding to the lower and upper bounds of the desired level of prediction interval:

$$L_\alpha = \begin{cases} \alpha(\hat{y}_t - y_t) & \text{if } \hat{y}_t \geq y_t \\ (1-\alpha)(y_t - \hat{y}_t) & \text{if } y_t > \hat{y}_t \end{cases}$$

Then $\mathrm{P}(\hat{y}_{t,lower} \leq y_t < \hat{y}_{t,upper}) = 1 - 2\alpha$.

4

Figure 3: τ controls the desired imbalance in the quantile forecast.

3.4 Training of the network

Most of the parameters set by Smyl in his initial model are preserved and result in good performances. The only changes we made was regarding the learning rate value and that we used a learning rate scheduler. All experiments for ES-RNN were made using Adam. We trained for 100 epochs starting with a learning rate of 0.01 which was progressively reduced to 1×10^{-5}. We also changed the batch size as detailed in the previous section and replaced the Tanh activation function with ReLu and LeakyReLu without noticeable improvements.

4 N-BEATS

4.1 Model Description

N-Beats architecture by design is simple and generic, yet expressive (deep). It does not rely on any time-series specific feature engineering or input scaling. It is also extendable and the basic components are:

- **Block**

  A block accepts an input x which is a history lookback window, also known as "backcast". And the block outputs two vectors $\hat{x}$ and $\hat{y}$, $\hat{x}$ is the block's best estimate of x, $\hat{y}$ is the block's forward forecast of length H. The length of the input window is a multiple of the forecast horizon H, and x value range is $2H$ to $7H$. Each block removes the portion of the signal it can explain well, making the forecast job of the downstream blocks easier. This also facilitates more fluid gradient backpropagation.

- **Stack**

  A stack is a collection of blocks and stacks can be hierarchically organized to make a very deep architecture. Each block provides a partial forecast that is first aggregated at the stack level and then at the overall network level.
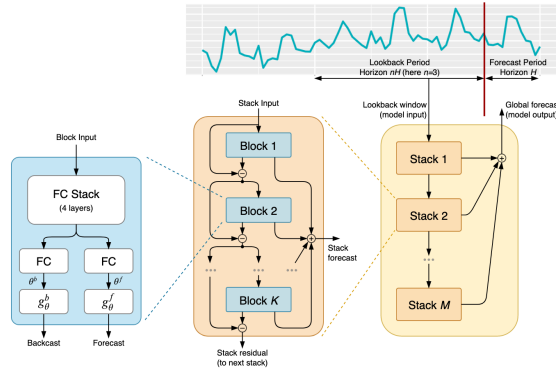

Figure 4: N-BEATS overall architecture.

References

References follow the acknowledgments. Use unnumbered first-level heading for the references. Any choice of citation style is acceptable as long as you are consistent. It is permissible to reduce the font size to `small` (9 point) when listing the references. **Remember that you can use more than eight pages as long as the additional pages contain *only* cited references.**

[1] Alexander, J.A. & Mozer, M.C. (1995) Template-based algorithms for connectionist rule extraction. In G. Tesauro, D.S. Touretzky and T.K. Leen (eds.), *Advances in Neural Information Processing Systems 7*, pp. 609–616. Cambridge, MA: MIT Press.

[2] Bower, J.M. & Beeman, D. (1995) *The Book of GENESIS: Exploring Realistic Neural Models with the GEneral NEural SImulation System.* New York: TELOS/Springer–Verlag.

[3] Hasselmo, M.E., Schnell, E. & Barkai, E. (1995) Dynamics of learning and recall at excitatory recurrent synapses and cholinergic modulation in rat hippocampal region CA3. *Journal of Neuroscience* **15**(7):5249-5262.